



MySQL InnoDB Cluster

Last Updated: January 2025

Copyright © 2023 Morpheus Data, LLC. All Rights Reserved
All third-party product and company names are property of their
respective holders and use does not imply any specific endorsement

Contents

Overview	3
Purpose	3
General	3
Components	3
MySQL	3
MySQL Shell	4
MySQL Router	4
Installation	4
MySQL	4
MySQL Shell	5
MySQL Router	5
MySQL Shell Commands.....	6
Start MySQL Shell.....	6
Exit MySQL Shell	6
Save History	6
Connect to a cluster node	6
Get the Clusters Status	6
Get the Clusters Extended Status	6
Get the Cluster Set Status (For Multi Site).....	6
Get the Cluster Set Extended Status (For Multi Site)	6
Check if DB Node is Ready for Configuration.	6
Configure DB Node to Meet Requirements for Clustering.....	7
Create a Single Site Cluster.....	7
Create a Multi Site Cluster.....	7
Remove a Node from the Cluster	7
Add a Node to the Cluster	7
Run a Script.....	7
Planned Failover to Another Cluster Member at the Same Site.....	7
Planned Failover to Another Site	7
All Nodes Down Recovery	8
Emergency Failover When a Site is Down	8
Emergency Failover Recovery of Down site.	8

Upgrades.....	8
Upgrade Order	8
Upgrade MySQL Router.....	8
Upgrade MySQL Shell.....	8
Upgrade MySQL	9
Post Upgrade Tasks	9
Metadata Upgrade.....	9
Upgrade Morpheus in a Multi-Site Environment	9
Morpheus Appliance Backup Settings	10
MySQL settings for MySQL Dump on Morpheus App Nodes	10
Create the following file on the Morpheus App nodes. This will set the proper setting for a MySQL dump db backup using InnoDB Cluster.	10

Overview

Purpose

The purpose of this document is to provide information on supporting MySQL InnoDB Cluster and its components. This document will cover the components that make up the solution, installation of the components and commands to set up, support and troubleshoot the environment.

General

The purpose of this document is to provide information on supporting MySQL InnoDB Cluster and its components. This document will cover the components that make up the solution, installation of the components and commands to setup, support and troubleshoot the environment.

Components

MySQL

The MySQL software delivers a very fast, multithreaded, multi-user, and robust SQL (Structured Query Language) database server. MySQL Server is intended for mission-critical, heavy-load production systems as well as for embedding into mass-deployed software. MySQL will be the foundational relational database used by Morpheus.

For additional information on MySQL 8:

<https://dev.mysql.com/doc/refman/8.0/en/introduction.html>

MySQL Shell

MySQL Shell is an advanced client and code editor for MySQL. In addition to the provided SQL functionality, MySQL Shell provides scripting capabilities for JavaScript and Python and includes APIs for working with MySQL. X DevAPI enables you to work with both relational and document data. AdminAPI enables you to work with InnoDB Cluster, InnoDB ClusterSet, and InnoDB ReplicaSet.

For additional information on MySQL Shell:
<https://dev.mysql.com/doc/mysql-shell/8.0/en/>

MySQL Router

MySQL Router is lightweight middleware that provides transparent routing between your application and any backend MySQL Servers. This router can run directly on the Morpheus application nodes and be bootstrapped from the MySQL cluster.

For additional information on MySQL Router:
<https://dev.mysql.com/doc/mysql-router/8.0/en/>

Installation

MySQL

You can install MySQL from the OS repositories or add the MySQL community repository. There are no additional 3rd party installations required for an InnoDB Cluster.

Using OS repositories:

Check for available versions from the repo.

```
apt - apt-cache show mysql-server
yum - yum list mysql-server
```

If the version is 8.0.32 or greater then install.

```
apt - apt-get update
      - apt-get -y install mysql-server
yum - yum update
      - yum -y install mysql-server
```

Using the MySQL Repos.

- yum - <https://dev.mysql.com/doc/refman/8.0/en/linux-installation-yum-repo.html>
- apt - <https://dev.mysql.com/doc/refman/8.0/en/linux-installation-apt-repo.html>

Configuration – After installation you will need to configure a few settings for use with the cluster.

Set the root password and create an admin account you will use as your Cluster Admin.

- ALTER USER 'root'@'localhost' IDENTIFIED WITH caching_sha2_password BY 'password';
- CREATE USER 'clusterAdmin'@'%' IDENTIFIED BY 'password';
- GRANT ALL PRIVILEGES ON *.* TO 'clusterAdmin'@'%' with grant option;

Set to generate invisible primary keys (InnoDB Cluster requires all tables to have primary keys. This is required to allow Morpheus to run on an InnoDB Cluster.

- set persist sql_generate_invisible_primary_key=1;

Add the following to the my.cnf

- innodb_buffer_pool_size=6G #This should be set to a size of 70 – 80% of the installed RAM
- innodb_buffer_pool_instances=6 #This should be set to one instance per gig of the pool size
- innodb_use_fdatasync=ON # This changes the sync method for better performance.
- bind-address=0.0.0.0 # This is binding MySQL to all interfaces.

Restart the MySQL service.

- systemctl restart mysql or systemctl restart mysqld

MySQL Shell

MySQL Shell can be installed on any machine you want to manage MySQL from. We will typically install this on each Morpheus app node to have it available for setup and management during the install and post deployment.

Manual Install

- go to <https://dev.mysql.com/downloads/shell/> and download the correct installer for your Linux distro.
- Install the package.

From repo

- Follow these steps to add the repo. <https://dev.mysql.com/doc/mysql-shell/8.0/en/mysql-shell-install-linux-quick.html>
- Then use apt-get install mysql-shell or yum install mysql-shell

MySQL Router

MySQL Router will be installed on each Morpheus App node. This will be the middleware that monitors the DB Cluster and points the Morpheus app node to the correct Primary DB node. Morpheus will connect to it on 127.0.0.1 and the R/W default port of 6446.

Manual Install

- go to <https://dev.mysql.com/downloads/router/> and download the correct installer for your Linux distro.
- Install the package.

From Repo

- Follow these steps to add the repo. <https://dev.mysql.com/doc/mysql-router/8.0/en/mysql-router-installation-linux.html>

- Then use apt-get install mysql-router or yum install mysql-router

Configuration

- Run the following to bootstrap the InnoDBCluster. With this config MySQL Router will get its config from the cluster and update based on the primary node that is active across any site in the cluster set.
 - o `mysqlrouter --bootstrap clusterAdmin@adb-5:3306 --account routeruser --user=mysqlrouter --disable-rest`
- Restart the router for changes to take effect.
 - o `systemctl restart mysqlrouter`
- Confirm the router is listening on the ports.
 - o `sudo lsof -i -P -n | grep LISTEN | grep mysqlrout`

MySQL Shell Commands

Start MySQL Shell

`mysqlsh`

Exit MySQL Shell

`\exit`

Save History

`\option --persist history.autoSave=1`

Connect to a cluster node

`\c clusterAdmin@dba-1:3306` #dba-1 is the db server name in this example

Get the Clusters Status

`\c clusterAdmin@dba-1:3306` #dba-1 is the db server name in this example
`cluster = dba.getCluster()`
`cluster.status()`

Get the Clusters Extended Status

`\c clusterAdmin@dba-1:3306` #dba-1 is the db server name in this example
`cluster = dba.getCluster()`
`cluster.status({extended: 1})`

Get the Cluster Set Status (For Multi Site)

`\c clusterAdmin@dba-1:3306` #dba-1 is the db server name in this example
`clusterset = dba.getClusterSet()`
`clusterset.status()`

Get the Cluster Set Extended Status (For Multi Site)

`\c clusterAdmin@dba-1:3306` #dba-1 is the db server name in this example
`clusterset = dba.getClusterSet()`
`clusterset.status({extended: 1})`

Check if DB Node is Ready for Configuration.

`dba.checkInstanceConfiguration('clusterAdmin@dba-1:3306')`

Configure DB Node to Meet Requirements for Clustering

```
dba.configureInstance('clusterAdmin@dba-1:3306')
```

Create a Single Site Cluster

Make sure you have run the “Configure DB Node to Meet Requirements” on each of the Cluster nodes first.

```
\c clusterAdmin@dba-1:3306      #Connect to one of the cluster nodes
cluster = dba.createCluster("A") # Create the cluster. "A" is the cluster name in this example
cluster.addInstance("dba-2:3306") #Add the second node to the cluster.
cluster.addInstance("dba-3:3306") #Add the third node to the cluster.
cluster.status()                #Get the cluster status to confirm functional
```

Create a Multi Site Cluster

Make sure you have run the “Configure DB Node to Meet Requirements” on each of the Cluster nodes first.

```
\c clusterAdmin@dba-1:3306      #Connect to one of the cluster nodes
cluster = dba.createCluster("A") # Create the cluster. "A" is the cluster name in this example
cluster.addInstance("dba-2:3306") #Add the second node to the cluster.
cluster.addInstance("dba-3:3306") #Add the third node to the cluster.
clusterset = cluster.createClusterSet("ClusterSet") #Create the clusterset "ClusterSet"
clusterb = clusterset.createReplicaCluster("dbb-1:3306", "B") #Create the replica cluster "B"
clusterb.addInstance("dbb-2:3306") #Add the second node to Cluster "B"
clusterb.addInstance("dbb-3:3306") #Add the third node to Cluster "B"
clusterset.status()             #Get the cluster set status and confirm functionality
```

Remove a Node from the Cluster

```
\c clusterAdmin@dba-1:3306      #Connect to one of the cluster nodes
cluster = dba.getCluster()      # Get the cluster
cluster.removeInstance('clusterAdmin@dba-2:3306') #We are removing the node dba-2
cluster.status()                # confirm status
```

Add a Node to the Cluster

```
\c clusterAdmin@dba-1:3306      #Connect to one of the cluster nodes
cluster = dba.getCluster()      # Get the cluster
cluster.addInstance('clusterAdmin@dba-2:3306') #We are adding the node dba-2
cluster.status()                # confirm status
```

Run a Script

```
mysqlsh --file myscript.js
```

Planned Failover to Another Cluster Member at the Same Site

```
\c clusterAdmin@dba-1:3306      #Connect to one of the cluster nodes
cluster = dba.getCluster()      # Get the cluster.
cluster.setPrimaryInstance('dba-2:3306') #We are setting the new primary member
cluster.status()                # confirm status
```

Planned Failover to Another Site

```
\c clusterAdmin@dba-1:3306      #Connect to one of the cluster nodes
cs = dba.getClusterSet()        # Get the Cluster Set.
cs.setPrimaryCluster("B")       #B is our second site in this example
cs.status()                     # confirm status
```

All Nodes Down Recovery

```
\c clusterAdmin@dba-1:3306      #Connect to one of the cluster nodes
dba.rebootClusterFromCompleteOutage()
```

Emergency Failover When a Site is Down

This should only be done if Primary site cannot be brought up first. Data loss could occur from Point in time of last replication.

```
\c clusterAdmin@dbb-1:3306      #Connect to one of the cluster nodes at the failover site.
clusterset = dba.getClusterSet() #Get Cluster Set.
clusterset.status()              #Check the Status. This will show you the site names.
clusterset.forcePrimaryCluster("B") #This will force failover to site "B".
```

Emergency Failover Recovery of Down site.

Once power is restored to Site A nodes. You can go through this repair process.

```
\c clusterAdmin@dba-1:3306      #Connect to one of the cluster nodes at the primary site.
dba.rebootClusterFromCompleteOutage() #Recover from all the nodes being down.
clusterset = dba.getClusterSet()      #Get the cluster Set.
clusterset.rejoinCluster("A")         #Rejoin the Cluster "A"
```

Upgrades

<https://dev.mysql.com/doc/mysql-shell/8.0/en/mysql-innodb-cluster-upgrade-rolling.html>

Upgrade Order

MySQL Router
MySQL Shell
MySQL Server

*If multi-site you will upgrade the replica (secondary) **sites** first. Then the primary site.

Upgrade MySQL Router

MySQL Router should be installed on each of your Morpheus App Nodes. You should upgrade one node at a time and confirm its up before moving to the next node in a rolling fashion.

- Download the package manually or use your added repo from the MySQL Router section above.
- `systemctl stop mysqlrouter`
- upgrade the package
- `systemctl start mysqlrouter`
- confirm MySQL Router is listening on 6446 and 6447

Upgrade MySQL Shell

It is best to always be running the latest version of MySQL Shell even if your version of MySQL is older. MySQL Shell should be installed on all Morpheus App Nodes.

- Download the package manually or use your added repo from the MySQL Shell section above.
- upgrade the package
- confirm you can run MySQL Shell `mysqlsh`

Upgrade MySQL

If multi-site you will upgrade the replica (secondary) sites first. Then the primary site.

Currently only 8.0.x versions of MySQL are supported versions such as 8.4.x are not supported.

MySQL Should be installed on at least 3 DB nodes per site external from the Morpheus App Nodes. You should upgrade the secondary **nodes** in the cluster first and then the Primary node at the site.

- Confirm which nodes are Secondary and Primary with mysqlsh
 - `\c clusterAdmin@dbb-1:3306` #Connect to a node
 - `cluster = dba.getCluster()`
 - `cluster.status()`
- Download the package manually or use your added repo from the MySQL section above.
- Stop the MySQL Service
 - `systemctl stop mysqld` or `systemctl stop mysql`
- Upgrade MySQL
- Start the MySQL Service
 - `systemctl start mysqld` or `systemctl start mysql`

Post Upgrade Tasks

- Check the cluster status after all upgrades
 - `\c clusterAdmin@dbb-1:3306` #Connect to a node
 - `cluster = dba.getCluster()`
 - `cluster.status({extended: true})` # get the extended status
 - Look to see if there are any errors listed for any of the nodes.
 - If any errors reported run the commands stated in the status to fix them.

Metadata Upgrade

After a MySQL upgrade you may need to upgrade the clusters metadata.

```
\c clusterAdmin@dba-1:3306      #Connect to one of the cluster nodes.
dba.upgradeMetadata()           #This will upgrade the metadata.
```

Upgrade Morpheus in a Multi-Site Environment

This guide assumes you are following the general recommended architecture for a multi-site deployment.

Other than upgrading Morpheus Application nodes, the morpheus-ui service should be down at the secondary site to make sure users are not accidentally using both sites simultaneously, as this will lead to issues.

•

To check if the morpheus-ui service is down on the application nodes you can run **morpheus-ctl status morpheus-ui** from the command line.

Environment

- Site A (Primary) and Site B (Secondary)
- 3 Application nodes and 3 Database nodes at each site
- morpheus-ui service is down at the secondary site.

Process

- Upgrade Site A following the standard HA environment upgrade process
- If required, conduct a failover of the database to Site B
- Upgrade Site B following the standard HA environment upgrade process
- Stop the morpheus-ui service on the app nodes at Site B if it is running
 - morpheus-ctl stop morpheus-ui
- If required, failback the database to Site A

Performing a database failover to the secondary site may be necessary in the upgrade process under certain circumstances, such as when latency is too high between the two sites or when application nodes at secondary site lack access to the primary sites database nodes.

Morpheus Appliance Backup Settings

MySQL settings for MySQL Dump on Morpheus App Nodes

Create the following file on the Morpheus App nodes. This will set the proper setting for a MySQL dump db backup using InnoDB Cluster.

/etc/my.cnf

Add the following config to it.

[mysqldump]

set-gtid-purged=OFF